

Ricardo Gemignani

Data Platform Consultant & Engineering Leader

Minneapolis, MN · ricardo@gemin.us · gemin.us · [LinkedIn](#)

SUMMARY

Data platform leader with 10+ years building and scaling data infrastructure at companies like Booking.com and Swiss Re. Specialize in pipeline optimization, data migrations, and building high-performing data teams. Now helping growing companies solve their toughest data infrastructure problems through Geminus.

EXPERIENCE

Founder & Principal Consultant

Geminus

Jan 2025 – Present · Minneapolis (Remote)

Data platform consulting practice helping companies optimize infrastructure, fix pipeline issues, and migrate to modern architectures.

- Delivered engagements across UK, Netherlands, and Germany: data collection infrastructure, pipeline optimization, data migration
- Built scalable data collection system for UK tech company enabling market intelligence capabilities
- Diagnosed and resolved critical pipeline failures for European fintech, restoring production reliability
- Designed migration strategy for German enterprise with comprehensive validation framework ensuring data integrity

Senior Data Engineer → Data Platform Lead

iptiQ (Swiss Re)

Jun 2021 – Dec 2024 · Amsterdam (Remote)

Led data platform engineering at Swiss Re's InsurTech subsidiary, scaling team from 4 to 20.

- Reduced critical pipeline runtime 96% (9h → 20min) architecting lakehouse migration (S3, Athena, Glue, dbt, Airflow)
- Scaled data team 4→20 as core technical interviewer, designing structured hiring framework
- Owned 3 business-critical domains managing 6+ engineers and 6-7 senior stakeholder relationships
- Designed cost-efficient architecture (serverless Athena vs Snowflake), 10x cheaper while maintaining performance
- Transformed engineering culture: code reviews, shared Python libraries, 12-week mentorship program
- Reduced incidents ~80% through proactive monitoring, data contracts, and stakeholder communication

AWS (S3, Athena, Glue, RDS, ECS, MWAA), Python, SQL, dbt, Airflow, Terraform

Core Developer - Data Engineering

[Booking.com](#)

Dec 2018 – Jun 2021 · Amsterdam

Data engineer on marketing platforms supporting ~\$12B annual ad spend.

- Operated petabyte-scale PPC pipelines (~1 PB/day) on 16-node bare-metal cluster, 99%+ SLA
- Maintained 200+ production Perl pipelines with strict 9-hour processing windows
- Built Kubernetes-native pipelines feeding Markov chain multi-touch attribution (trillions of events)
- Migrated ephemeral outputs to BigQuery, improving data access for analytics teams
- On-call "fort-holder" every 6 weeks, triaging critical failures affecting multi-million-euro decisions

Hadoop, BigQuery, Kafka, Kubernetes, Perl, Python

Senior Software Engineer → Team Lead (Machine Learning)

[Talpa Network / Emesa](#)

Sep 2016 – Dec 2018 · Amsterdam

Led ML engineering at [vakantieveilingen.nl](#), one of Netherlands' largest online auction platforms.

- Productionized auction pricing system processing ~800M auctions/month
- Built real-time prediction system serving ~300 auctions/sec (online regression, vanilla Python)
- Led two ML teams (auction optimization, recommendation engine), establishing ML engineering practices
- Founded Python guild (~10-15 engineers), creating shared libraries and coding standards
- Bridged data science and business, translating "profit per minute" into production systems

Python, Kafka, Avro, MySQL, Redis, Elasticsearch, Docker

Senior Software Engineer & Team Lead

[Grupo RBS](#)

Dec 2014 – Jul 2016 · Porto Alegre, Brazil

Led technical redesign of CMS at major Brazilian media company (8 newspapers, 2 TV channels, 8 radio stations).

- Reduced content publishing time from hours to minutes with microservices architecture
- Technical lead for ~6 engineers, guiding architecture and mentoring on service-oriented design
- Built weather data platform aggregating international + local Brazilian sources
- Introduced caching layers (Varnish, Redis) improving editorial workflow performance

Java, Python, Ruby, JavaScript, MongoDB, Oracle, Docker, Redis

Founder & Technical Director

[Fleye](#)

Feb 2006 – Aug 2014 · Porto Alegre, Brazil

Founded and led bootstrapped software company for 8+ years.

- Built engineering team from zero, 50/50 hands-on development and leadership
- Delivered multiple production products across Java, Python, PHP, MySQL, PostgreSQL, MongoDB
- Created engineering culture that outlived the company—team members later founded their own company together
- Evolved leadership approach from high-pressure to encouragement-based, transforming team performance

SKILLS

Data: Data Engineering, Data Pipelines, Data Architecture, Data Migration, ETL/ELT, Data Quality

Cloud & Infrastructure: AWS (S3, Athena, Glue, RDS, ECS, MWAA), Terraform, Kubernetes, Docker

Languages & Tools: Python, SQL, dbt, Apache Airflow, Kafka, BigQuery, PostgreSQL

Leadership: Team Building, Technical Hiring, Stakeholder Management, Mentorship, Cross-functional Collaboration

LOCATION

Based in Minneapolis, MN. US work authorization. Open to remote and hybrid opportunities.